

```
...
Purpose of the model building phase:
- Implementing the design decisions from the model planning phase
- Training the models across the defined hyperparameter configurations
- Evaluating each model using the chosen metrics (RMSE, MAE, MSE)
- Logging parameters, metrics, and model artifacts via MLflow for comparison and reuse
...
```

```
'\nPurpose of the model building phase:\n- Implementing the design decisions from the model planning phase\n- Training the models across the defined hyperparameter configurations\n- Evaluating each model using the chosen metrics (RMSE, MAE, MSE)\n- Logging parameters, metrics, and model artifacts via MLflow for comparison and reuse\n'
```

```
...
Model Building Summary/ Results of model building:
In the model building phase, the preprocessing and prediction steps were combined into a single, reproducible scikit-learn Pipeline consisting of a StandardScaler and a RandomForestRegressor. The data was split chronologically into train, validation and test sets (80% / 10% / 10%), as shuffling is not appropriate for time-dependent data. A baseline Random Forest model was trained and evaluated, followed by a hyperparameter tuning step using RandomizedSearchCV. The search was executed inside an MLflow-tracked training allowing parameters, metrics and artifacts to be logged and compared across experiments. The best configuration found within defined parameter ranges was: n_estimators = 200, min_samples_leaf = 3, max_features = "sqrt", max_depth = 10.
```

The tuned model showed only a small improvement over the baseline, with validation metrics RMSE = 6.20, MAE = 4.49 and MSE = 10. These values are the best combination sampled by randomized search but do not represent a guaranteed global optimum. The final best\_estimator\_ was logged to MLflow and will be used as the model selected for deployment.

```
...
```

```
'\nModel Building Summary/ Results of model building:\nIn the model building phase, the preprocessing and prediction steps were combined into a single, \nreproducible scikit-learn Pipeline consisting of a StandardScaler and a RandomForestRegressor. \nThe data was split chronologically into train, validation and test sets (80% / 10% / 10%), \nas shuffling is not appropriate for time-dependent data. A baseline Random Forest model was trained and evaluated, \nfollowed by a hyperparameter tuning step using RandomizedSearchCV. The search was executed inside an MLflow-tracked training run, \nallowing parameters, metrics and artifacts to be logged and compared across experiments. The best configuration found within the \nundefined parameter ranges was: n_estimators = 200, min_samples_leaf = 3, max_features = "sqrt", max_depth = 10. \n\nThe tuned model showed only a small improvement over the baseline, with validation metrics RMSE = 6.20, MAE = 4.49 and MSE = 38.49. \nThese values are the best combination sampled by randomized search but do not represent a guaranteed global optimum. \n\nThe final best_estimator_ was logged to MLflow and will be used as the model selected for deployment.\n\n'
```

```
import pandas as pd

!pip install scikit-learn
!pip install mlflow
!pip install cloudpickle
import mlflow
import numpy as np

from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestRegressor
from sklearn.pipeline import Pipeline
from sklearn.metrics import mean_squared_error, mean_absolute_error
```

```
Collecting scikit-learn
  Downloading scikit_learn-1.7.2-cp312-cp312-manylinux2014_x86_64.manylinux_2_17_x86_64.whl.metadata (11 kB)
Requirement already satisfied: numpy>=1.22.0 in /opt/conda/lib/python3.12/site-packages (from scikit-learn) (2.3.5)
Collecting scipy>=1.8.0 (from scikit-learn)
  Downloading scipy-1.16.3-cp312-cp312-manylinux2014_x86_64.manylinux_2_17_x86_64.whl.metadata (62 kB)
Collecting joblib>=1.2.0 (from scikit-learn)
  Downloading joblib-1.5.2-py3-none-any.whl.metadata (5.6 kB)
Collecting threadpoolctl>=3.1.0 (from scikit-learn)
  Downloading threadpoolctl-3.6.0-py3-none-any.whl.metadata (13 kB)
Download scikit_learn-1.7.2-cp312-cp312-manylinux2014_x86_64.manylinux_2_17_x86_64.whl (9.5 MB)
  9.5/9.5 MB 18.5 MB/s 0:00:00m0:00:0100:01
Downloading joblib-1.5.2-py3-none-any.whl (308 kB)
Download scikit_learn-1.7.2-cp312-cp312-manylinux2014_x86_64.manylinux_2_17_x86_64.whl (35.7 MB)
  35.7/35.7 MB 19.7 MB/s 0:00:01m0:00:0100:02
Downloading threadpoolctl-3.6.0-py3-none-any.whl (18 kB)
Installing collected packages: threadpoolctl, scipy, joblib, scikit-learn
  4/4 [scikit-learn][0m [scikit-learn]
Successfully installed joblib-1.5.2 scikit-learn-1.7.2 scipy-1.16.3 threadpoolctl-3.6.0
Collecting mlflow
  Downloading mlflow-3.6.0-py3-none-any.whl.metadata (31 kB)
Collecting mlflow-skinnny==3.6.0 (from mlflow)
  Downloading mlflow_skinnny-3.6.0-py3-none-any.whl.metadata (31 kB)
Collecting mlflow-tracing==3.6.0 (from mlflow)
  Downloading mlflow_tracing-3.6.0-py3-none-any.whl.metadata (19 kB)
Collecting Flask-CORS<7 (from mlflow)
  Downloading flask_cors-6.0.1-py3-none-any.whl.metadata (5.3 kB)
Collecting Flask<4 (from mlflow)
```

```

  Downloading flask-3.1.2-py3-none-any.whl.metadata (3.2 kB)
Requirement already satisfied: alembic!=1.10.0,<2 in /opt/conda/lib/python3.12/site-packages (from mlflow) (1.17.2)
Requirement already satisfied: cryptography<47,>=43.0.0 in /opt/conda/lib/python3.12/site-packages (from mlflow) (46.0.3)
Collecting docker<8,>=4.0.0 (from mlflow)
  Downloading docker-7.1.0-py3-none-any.whl.metadata (3.8 kB)
Collecting graphene<4 (from mlflow)
  Downloading graphene-3.4.3-py2.py3-none-any.whl.metadata (6.9 kB)
Collecting gunicorn<24 (from mlflow)
  Downloading gunicorn-23.0.0-py3-none-any.whl.metadata (4.4 kB)
Collecting huey<3,>=2.5.0 (from mlflow)
  Downloading huey-2.5.4-py3-none-any.whl.metadata (4.6 kB)
Requirement already satisfied: matplotlib<4 in /opt/conda/lib/python3.12/site-packages (from mlflow) (3.10.8)
Requirement already satisfied: numpy<3 in /opt/conda/lib/python3.12/site-packages (from mlflow) (2.3.5)
Requirement already satisfied: pandas<3 in /opt/conda/lib/python3.12/site-packages (from mlflow) (2.3.3)
Requirement already satisfied: pyarrow<23,>=4.0.0 in /opt/conda/lib/python3.12/site-packages (from mlflow) (21.0.0)
Requirement already satisfied: scikit-learn<2 in /opt/conda/lib/python3.12/site-packages (from mlflow) (1.7.2)
Requirement already satisfied: scipy<2 in /opt/conda/lib/python3.12/site-packages (from mlflow) (1.16.3)
Requirement already satisfied: sqlalchemy<3,>=1.4.0 in /opt/conda/lib/python3.12/site-packages (from mlflow) (2.0.44)
Requirement already satisfied: cachetools<7,>=5.0.0 in /opt/conda/lib/python3.12/site-packages (from mlflow-skinny==3.6.0->mlflow)
Requirement already satisfied: click<9,>=7.0 in /opt/conda/lib/python3.12/site-packages (from mlflow-skinny==3.6.0->mlflow)
Requirement already satisfied: cloudpickle<4 in /opt/conda/lib/python3.12/site-packages (from mlflow-skinny==3.6.0->mlflow)
Collecting databricks-sdk<1,>=0.20.0 (from mlflow-skinny==3.6.0->mlflow)
  Downloading databricks_sdk-0.73.0-py3-none-any.whl.metadata (40 kB)
Collecting fastapi<1 (from mlflow-skinny==3.6.0->mlflow)
  Downloading fastapi-0.123.8-py3-none-any.whl.metadata (30 kB)
Requirement already satisfied: gitpython<4,>=3.1.9 in /opt/conda/lib/python3.12/site-packages (from mlflow-skinny==3.6.0->mlflow)
Requirement already satisfied: importlib_metadata!=4.7.0,<9,>=3.7.0 in /opt/conda/lib/python3.12/site-packages (from mlflow)
Collecting opentelemetry-api<3,>=1.9.0 (from mlflow-skinny==3.6.0->mlflow)
  Downloading opentelemetry_api-1.39.0-py3-none-any.whl.metadata (1.5 kB)
Collecting opentelemetryproto<3,>=1.9.0 (from mlflow-skinny==3.6.0->mlflow)

```

```

import os
os.listdir()

```

```

['mlruns',
 '.ipynb_checkpoints',
 'file_modelbuilding.py',
 'file_dis&dataprep.py',
 'Model Building.ipynb',
 'file_modelbuilding_uv.py',
 'Discovery&DataPrep.ipynb',
 'Model_planning.ipynb']

```

```

#loading df_merged and splitting data
df_merged = pd.read_csv("/work/bda2/datasets/df_merged.csv")
#1:1 the same function as in discovery&dataprep notebook, but I have to define it for each notebook
def create_splits(df_merged, seed=42):

    df_merged = df_merged.sort_values(by="target_time").reset_index(drop=True)

    n = len(df_merged)

    split_1 = int(n * 0.80)
    temp_train = df_merged.iloc[:split_1]
    test_df = df_merged.iloc[split_1:]

    # 2 Train/Validation split (80/20 of temp_train)
    split_2 = int(len(temp_train) * 0.80)
    train_df = temp_train.iloc[:split_2]
    val_df = temp_train.iloc[split_2:]

    return train_df, val_df, test_df

# create the splits
train_df, val_df, test_df = create_splits(df_merged, seed=42)

print(len(train_df), len(val_df), len(test_df))

#appying general function def create_splits to df_merged
train_df, val_df, test_df = create_splits(df_merged)

```

```

169 43 53

```

```

df_merged.head()

```

	target_time	Direction	Lead_hours	Source_time	Speed	Total	dir_deg	dir_rad	u	v
0	2021-12-11 15:00:00+00:00	SSE	1	1639227600	11.17600	27.621048	157.5	2.748894	-10.325278	4.276870e+00
1	2021-12-11 18:00:00+00:00	SSW	1	1639238400	8.04672	21.135542	202.5	3.534292	-7.434200	-3.079346e+00
2	2021-12-11 21:00:00+00:00	WSW	1	1639249200	11.17600	20.616209	247.5	4.319690	-4.276870	-1.032528e+01
...	2021-12-12	...	...	...	...	...	...	...	...	...

```
#defining features and label (supervised learning)
numeric_features = ["Speed", "u", "v"] #needed for preprocessing, i.e., the preprocessor knows which features to scale
feature_cols = ["Speed", "u", "v"]
keep_cols = ["Speed", "u", "v", "Total"]
#only keeping the nesscar columns. I also decided to not keep target_time, since not time series analysis
train_df = train_df[keep_cols]
val_df = val_df[keep_cols]
test_df = test_df[keep_cols]
X_train = train_df[feature_cols]
y_train = train_df["Total"]

X_val = val_df[feature_cols]
y_val = val_df["Total"]
```

```
#defining evaluation function
def evaluate_model(model, X_val, y_val, name="model"):
    preds = model.predict(X_val)

    rmse = np.sqrt(mean_squared_error(y_val, preds))
    mae = mean_absolute_error(y_val, preds)
    mse = mean_squared_error(y_val, preds)

    print(f"\nResults for {name}:")
    print(f"RMSE: {rmse:.4f}")
    print(f"MAE : {mae:.4f}")
    print(f"MSE : {mse:.4f}")

    return rmse, mae, mse
```

```
# preprocessing and pipeline generic / modular approach
# preprocessing like in discovery & data prep
preprocessor = ColumnTransformer(
    [("num", StandardScaler(), numeric_features)],
    remainder="passthrough"
)

from sklearn.pipeline import Pipeline

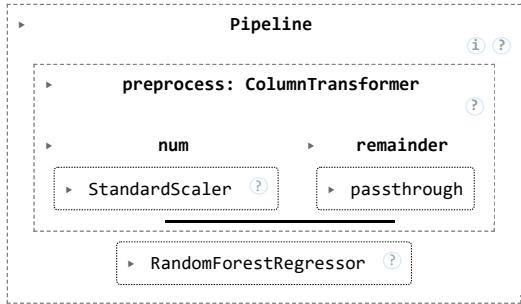
# generic pipeline builder
def make_pipeline(model):
    return Pipeline([
        ("preprocess", preprocessor),
        ("model", model)
    ])

...
Note: Although splitting and preprocessing were already tested in the Data Preparation notebook,# they must be implemented here in the Model Building phase. The previous steps were exploratory, this notebook performs the final, reproducible train/test splits and constructs the full preprocessing+model pipeline that will be trained, evaluated, and logged with MLflow.
...
'\nNote: Although splitting and preprocessing were already tested in the Data Preparation notebook,# they must be implemented again\nhere in the Model Building phase. The previous steps were exploratory, this notebook performs the final, reproducible train/validation/\ntest splits and constructs the full preprocessing+model pipeline that will be trained, evaluated, and logged with MLflow.\n'
```

```
#pipeline specif part for random forest, will be used for train/val/ test seperately to prevent data leakage
rfr_baseline = RandomForestRegressor(
    n_estimators=300,
    max_depth=10,
    min_samples_leaf=3,
    max_features="sqrt",
    random_state=42
)
```

```
pipeline_rfr_baseline = make_pipeline(rfr_baseline)
```

```
#training the model
pipeline_rfr_baseline.fit(X_train, y_train)
```



```
rmse_rfr, mae_rfr, mse_rfr = evaluate_model(
    pipeline_rfr_baseline, X_val, y_val, "RandomForest"
)
```

Results for RandomForest:

RMSE: 6.2040  
MAE : 4.4881  
MSE : 38.4901

```
# Collect results for a clean model comparison
results = {}

results["RandomForest"] = (rmse_rfr, mae_rfr, mse_rfr)
```

```
#for having context to evaluate results
df_merged["Total"].describe()
```

count	265.000000
mean	21.449641
std	10.795674
min	1.096348
25%	11.408295
50%	26.096300
75%	30.348545
max	35.492622
Name:	Total, dtype: float64

#The baseline Random Forest achieves an MAE of 4.49 MW (=21% of mean production) and an RMSE of 6.20 MW,  
#indicating that the model captures the general wind-power relationship but still shows occasional larger deviations due to

#need to optimize the model in another run. I use randomized search over grid search, since this hyperparameter tuning approach  
#is supposed to be better for small datasets.

```
import mlflow
import mlflow.sklearn
from sklearn.model_selection import RandomizedSearchCV

# Hyperparameter search space
param_grid = {
    "model__n_estimators": [200, 300, 500],
    "model__max_depth": [8, 10, 12, None],
    "model__min_samples_leaf": [1, 2, 3, 5],
    "model__max_features": ["sqrt", "auto"],
}

# RandomizedSearchCV setup
search = RandomizedSearchCV(
    estimator=pipeline_rfr_baseline,
    param_distributions=param_grid,
    n_iter=12,
    scoring="neg_mean_absolute_error",
    cv=3,
    n_jobs=-1,
    random_state=42
)
```

```
# --- MLflow Run ---
with mlflow.start_run(run_name="RandomForest_Finetuning"):

    search.fit(X_train, y_train)
    best_model = search.best_estimator_

    val_preds = best_model.predict(X_val)

    rmse_val = np.sqrt(mean_squared_error(y_val, val_preds))
    mae_val = mean_absolute_error(y_val, val_preds)
    mse_val = mean_squared_error(y_val, val_preds)

    # Log hyperparameters
    for param_name, param_value in search.best_params_.items():
        mlflow.log_param(param_name, param_value)

    # Log metrics
    mlflow.log_metric("rmse_val", rmse_val)
    mlflow.log_metric("mae_val", mae_val)
    mlflow.log_metric("mse_val", mse_val)

    # Log model
    mlflow.sklearn.log_model(best_model, "model")

    print("Best hyperparameters:", search.best_params_)
    print("MAE:", mae_val)
    print("RMSE:", rmse_val)
    print("MSE:", mse_val)
```

```
/opt/conda/lib/python3.12/site-packages/mlflow/tracking/_tracking_service/utils.py:140: FutureWarning: Filesystem tracking t
    return FileStore(store_uri, store_uri)
/opt/conda/lib/python3.12/site-packages/sklearn/model_selection/_validation.py:516: FitFailedWarning:
12 fits failed out of a total of 36.
The score on these train-test partitions for these parameters will be set to nan.
If these failures are not expected, you can try to debug them by setting error_score='raise'.
```

Below are more details about the failures:

```
-----  
7 fits failed with the following error:  
Traceback (most recent call last):  
  File "/opt/conda/lib/python3.12/site-packages/sklearn/model_selection/_validation.py", line 859, in _fit_and_score  
    estimator.fit(X_train, y_train, **fit_params)  
  File "/opt/conda/lib/python3.12/site-packages/sklearn/base.py", line 1365, in wrapper  
    return fit_method(estimator, *args, **kwargs)  
          ^^^^^^^^^^^^^^^^^^^^^  
  File "/opt/conda/lib/python3.12/site-packages/sklearn/pipeline.py", line 663, in fit  
    self._final_estimator.fit(Xt, y, **last_step_params["fit"])  
  File "/opt/conda/lib/python3.12/site-packages/sklearn/base.py", line 1358, in wrapper  
    estimator._validate_params()  
  File "/opt/conda/lib/python3.12/site-packages/sklearn/base.py", line 471, in _validate_params  
    validate_parameter_constraints()  
  File "/opt/conda/lib/python3.12/site-packages/sklearn/utils/_param_validation.py", line 98, in validate_parameter_constraints  
    raise InvalidParameterError(  
sklearn.utils._param_validation.InvalidParameterError: The 'max_features' parameter of RandomForestRegressor must be an int  
  
-----  
5 fits failed with the following error:  
Traceback (most recent call last):  
  File "/opt/conda/lib/python3.12/site-packages/sklearn/model_selection/_validation.py", line 859, in _fit_and_score  
    estimator.fit(X_train, y_train, **fit_params)  
  File "/opt/conda/lib/python3.12/site-packages/sklearn/base.py", line 1365, in wrapper  
    return fit_method(estimator, *args, **kwargs)  
          ^^^^^^^^^  
  File "/opt/conda/lib/python3.12/site-packages/sklearn/pipeline.py", line 663, in fit  
    self._final_estimator.fit(Xt, y, **last_step_params["fit"])  
  File "/opt/conda/lib/python3.12/site-packages/sklearn/base.py", line 1358, in wrapper  
    estimator._validate_params()  
  File "/opt/conda/lib/python3.12/site-packages/sklearn/base.py", line 471, in _validate_params  
    validate_parameter_constraints()  
  File "/opt/conda/lib/python3.12/site-packages/sklearn/utils/_param_validation.py", line 98, in validate_parameter_constraints  
    raise InvalidParameterError(  
sklearn.utils._param_validation.InvalidParameterError: The 'max_features' parameter of RandomForestRegressor must be an int  
  
    warnings.warn(some_fits_failed_message, FitFailedWarning)  
/opt/conda/lib/python3.12/site-packages/sklearn/model_selection/_search.py:1135: UserWarning: One or more of the test scores  
      nan           nan -4.87945125 -4.82538025           nan -4.65675384]  
    warnings.warn(  
2025/12/04 18:58:21 WARNING mlflow.models.model: `artifact_path` is deprecated. Please use `name` instead.  
2025/12/04 18:58:30 WARNING mlflow.models.model: Model logged without a signature and input example. Please set `input_examp  
Best hyperparameters: {'model_n_estimators': 200, 'model_min_samples_leaf': 3, 'model_max_features': 'sqrt', 'model_max_  
MAE: 4.456784194494066  
RMSE: 6.1691309172878  
MSE: 38.05816877155546
```

```
# Save tuned Random Forest results in dictionary
results["RandomForest_Tuned"] = (rmse_val, mae_val, mse_val)

...
evaluation metrics of run1 for comparison: RMSE (Root Mean Squared Error): 6.2040, MAE (Mean Absolute Error): 4.4881, MSE (Mean Squared Error): 38.4901

--> "only" a very small improvement
interpretation of best hyperparameters according to randomized search. for the given range defined in the code in the function parameter_grid, this values for the hyperparameters are the best combination:
'model__n_estimators': 200, 'model__min_samples_leaf': 3, 'model__max_features': 'sqrt', 'model__max_depth': 10

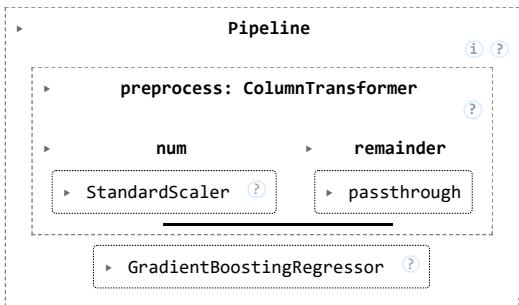
It is worth mentioning, that it is not by default the best combination possible. It is the best combination this hyperparameter optimization algorithm has found. Other hyperparameters might have come to other, "better" performing hyperparameter combinations.\n\n'
```

```
#using now gradient descent regressor to be able to compare the model performance
```

```
#using same generic data pipeline and preprocessing
#defining specific pipeline part for gradient descent regressor
from sklearn.ensemble import GradientBoostingRegressor

gbr_model = GradientBoostingRegressor(
    n_estimators=300,
    learning_rate=0.1,
    max_depth=3,
    min_samples_split=2,
    loss="squared_error",
    random_state=42
)
```

```
#applying pipeline
pipeline_gbr = make_pipeline(gbr_model)
#training model
pipeline_gbr.fit(X_train, y_train)
```



```
#evaluating the model with generic function defined above
rmse_gbr, mae_gbr, mse_gbr = evaluate_model(
    pipeline_gbr, X_val, y_val, "GradientBoosting"
)
```

```
Results for GradientBoosting:
RMSE: 7.7969
MAE : 5.1331
MSE : 60.7921
```

```
results["GradientBoosting"] = (rmse_gbr, mae_gbr, mse_gbr)
print(results)
```

```
{'RandomForest': (np.float64(6.20403528759331), 4.488051380047191, 38.49005384970301), 'RandomForest_Tuned': (np.float64(6.1
```

```
#in baseline comparison random forest clearly better
```

```
#optimizing gradientboosting
#same method as in random forest,i.e., randomized search space
from sklearn.model_selection import RandomizedSearchCV
from sklearn.ensemble import GradientBoostingRegressor
import mlflow
import mlflow.sklearn

# Gradient Boosting Hyperparameter Grid
param_grid_gbr = {
    "model__n_estimators": [100, 300, 500],
    "model__learning_rate": [0.1, 0.05, 0.01],
    "model__max_depth": [2, 3, 4],
    "model__min_samples_split": [2, 5, 10],
    "model__loss": ["squared_error"]
}
```

```
search_gbr = RandomizedSearchCV(
    estimator=gradient_boosting,
    param_distributions=param_grid_gbr,
    n_iter=12,
    scoring="neg_mean_absolute_error",
    cv=3,
    n_jobs=-1,
    random_state=42
)
```

```
with mlflow.start_run(run_name="GradientBoosting_Finetuning"):

    search_gbr.fit(X_train, y_train)
    best_gbr = search_gbr.best_estimator_

    rmse_val, mae_val, mse_val = evaluate_model(
        best_gbr, X_val, y_val, name="GradientBoosting_Tuned"
    )

    # Log hyperparameters
    for param, value in search_gbr.best_params_.items():
        mlflow.log_param(param, value)

    # Log metrics
    mlflow.log_metric("rmse_val", rmse_val)
    mlflow.log_metric("mae_val", mae_val)
    mlflow.log_metric("mse_val", mse_val)

    # Log fitted model as MLflow artifact
    mlflow.sklearn.log_model(best_gbr, "model")

    print("Best hyperparameters:", search_gbr.best_params_)
    print("MAE:", mae_val)
    print("RMSE:", rmse_val)
    print("MSE:", mse_val)
```

2025/12/04 18:58:31 WARNING mlflow.models.model: `artifact\_path` is deprecated. Please use `name` instead.

Results for GradientBoosting\_Tuned:

RMSE: 6.8582

MAE : 4.9408

MSE : 47.0356

2025/12/04 18:58:37 WARNING mlflow.models.model: Model logged without a signature and input example. Please set `input\_example`

Best hyperparameters: {'model\_\_n\_estimators': 100, 'model\_\_min\_samples\_split': 5, 'model\_\_max\_depth': 2, 'model\_\_loss': 'sq

MAE: 4.94079210935536

RMSE: 6.858247887874707

MSE: 47.03556409153787

```
results["GradientBoosting_Tuned"] = (rmse_val, mae_val, mse_val)
results
```

```
{'RandomForest': (np.float64(6.20403528759331),
 4.488051380047191,
 38.49005384970301),
'RandomForest_Tuned': (np.float64(6.169130309172878),
 4.456784194494066,
 38.05816877155546),
'GradientBoosting': (np.float64(7.7969309773830915),
 5.133091161779405,
 60.79213266607606),
```

```
'GradientBoosting_Tuned': (np.float64(6.858247887874707),
 4.940792109355536,
 47.03556409153787)}
```

```
#randomforest_tuned is the best model
```

```
#showing me the run id for the best model, so I can deploy it with mlflow
with mlflow.start_run(run_name="RandomForest_Tuned") as run:
    run_id = run.info.run_id
    print("Run ID:", run_id)
```

```
Run ID: ddb8e93ba4924decbcea9d58108a5241
```

```
#now applying the best model to future_df
# --- Prepare df_future the same way as df_merged ---

# Convert direction to degrees
df_future = pd.read_csv("../datasets/future.csv")
dir_map = {
    "N": 0, "NNE": 22.5, "NE": 45, "ENE": 67.5,
    "E": 90, "ESE": 112.5, "SE": 135, "SSE": 157.5,
    "S": 180, "SSW": 202.5, "SW": 225, "WSW": 247.5,
    "W": 270, "WNW": 292.5, "NW": 315, "NNW": 337.5
}

df_future["dir_deg"] = df_future["Direction"].map(dir_map)
df_future["dir_rad"] = np.deg2rad(df_future["dir_deg"])

# calculate u and v
df_future["u"] = df_future["Speed"] * np.cos(df_future["dir_rad"])
df_future["v"] = df_future["Speed"] * np.sin(df_future["dir_rad"])
```

```
#select features
X_future = df_future[["Speed", "u", "v"]]
```

```
future_predictions = best_model.predict(X_future)

df_future["Predicted_Power"] = future_predictions

df_future[["time", "Predicted_Power"]].head()
```

	time	Predicted_Power
0	2022-03-11 15:00:00+00:00	30.620526
1	2022-03-11 18:00:00+00:00	30.794534
2	2022-03-11 21:00:00+00:00	30.725303
3	2022-03-12 00:00:00+00:00	31.407954
4	2022-03-12 03:00:00+00:00	30.637139

```
#saving it
"""
df_future.to_csv("../datasets/future_predictions.csv", index=False)
"""

'\nndf_future.to_csv("../datasets/future_predictions.csv", index=False)\n'
```

```
#visualising the results

# Load predictions
df_future = pd.read_csv("../datasets/future_predictions.csv")

# Ensure time column is datetime
df_future["time"] = pd.to_datetime(df_future["time"])

# Sort by time (just to be safe)
df_future = df_future.sort_values("time")

# Plot
plt.figure(figsize=(12,6))
```

```
plt.plot(df_future["time"], df_future["Predicted_Power"], marker="o", linestyle="-")

plt.title("Predicted Future Power Output (3-hour intervals)")
plt.xlabel("Time")
plt.ylabel("Predicted Power (MW)")

# Format date labels
plt.gca().xaxis.set_major_formatter(mdates.DateFormatter('%Y-%m-%d\n%H:%M'))
plt.gca().xaxis.set_major_locator(mdates.AutoDateLocator())

plt.grid(True)
plt.tight_layout()
plt.show()
```

